

Working with NoSQL



Filip Ekberg

Principal Consultant & CEO

@fekberg fekberg.com






Non-relational Database

Commonly referred to as NoSQL

A group of many different data stores


NoSQL Databases



Key-value store or Key-value cache



Document store or Document database



Object storage





Table storage



Graph database




NoSQL Databases in This Module



Document database

JSON serialized data



Key-value store

Persistent dictionary




Table store


Tabular data, often terabytes of it




NoSQL or Relational Database



A relational database engine enforces relationships



Relational data is not verified by the NoSQL database engine



Data in NoSQL is accessed through keys

Can be used to define a relationship but this will **NOT** be enforced



Document Database

Store documents

JSON, XML, or other supported formats

All data in one place

Useful to store entities with all its related data in one place



Example: Document in a Document Database

```
{  
  "InvoiceId": "IN123",  
  "CustomerId": "1",  
  "Due": "2025-01-01",  
  "Items": [],  
  "Total": 1999.50  
}
```



Invoice

```
public class Invoice
{
    public string InvoiceId { get; set; }
    public Guid CustomerId { get; set; }
    public DateOnly Due { get; set; }
    public IEnumerable<Item> Items { get; set; }
    public decimal Total { get; set; }
}
```



Access a Document

Partition Key

A good partitioning leads to better scaling

Row Key

The key for the row in the specified partition



NoSQL databases are **often**
very **fast** and **scalable**



Key Value Store

Dictionary

JSON, binary, simple or complex types

Good for caching!

It's used like a hosted, scalable dictionary to cache data



Redis is a common key value
store



Feature Toggle

```
// Ask the Key Value Store for a Boolean value
if(FeatureService.IsEnabled("Order_Page_2.0"))
{
    ...
}
```



Table Store

Tabular data

Looks like any column from a traditional database, without relationships

Good for LOTS of data

Not uncommon to be used for terabytes of data




Document Database Content



Partition Key

Use to determine in which partition to store the data.
A good partition key makes scaling easier and better.



Row Key

Commonly referred to as id and is used to index the data in a partition.
Has to be unique!



Data

The data serialized to a supported format. Commonly stored as JSON.



Distributed Cache (Key Value Store)

```
public string GetDetailsFor(string itemId)
{
    var details = database.StringGet(itemId);

    if (details.HasValue)
    {
        return details;
    }
    var loadedDetails = ... ;

    database.StringSet(itemId, loadedDetails);

    return loadedDetails;
}
```



**Always access data through
its key!**

Not meant to be used to
iterate through all cached
items.



Introduce a (distributed)
cache to store data that is
frequently accessed



Redis is an in-memory cache


If the instance/server restart all
data is lost.



Using the Distributed Cache (**Key Value Store**)



Available in the cache?



Load the data a



Data

The data serialized to a supported format. Commonly stored as JSON.



Using a Table Store

```
TableServiceClient tableServiceClient = new("CONNECTION STRING");  
  
TableClient tableClient =  
    tableServiceClient.GetTableClient("SupportRequests");  
  
tableClient.CreateIfNotExists();  
  
tableClient.AddEntity(...);
```



Access a Row in the Table

Partition Key

A good partitioning leads to better scaling

Row Key

The key for the row in the specified partition




Choose a partition key that
evenly distributes your data


Warehouse identifier?
Customer Identifier?




Using More Than One Database



Using More Than One Database



Using More Than One Database



Cloud Hosted NoSQL Databases

Scalable

High performance



Know your pricing!

Always read the
documentation to avoid
unnecessary cost



Create a Document

```
using var client = new CosmosClient(  
    accountEndpoint: "ADD THE ENDPOINT HERE",  
    authKeyOrResourceToken: "ADD THE TOKEN HERE"  
);  
  
Database database =  
    await client.CreateDatabaseIfNotExistsAsync("WarehouseManagementSystem");  
  
Container container = await database.CreateContainerIfNotExistsAsync(  
    id: "invoices",  
    partitionKeyPath: "/customerId", throughput: 400);  
  
await container.UpsertItemAsync(invoice);
```



Create a Document

```
using var client = new CosmosClient(  
    accountEndpoint: "ADD THE ENDPOINT HERE",  
    authKeyOrResourceToken: "ADD THE TOKEN HERE"  
);  
  
Database database =  
    await client.CreateDatabaseIfNotExistsAsync("WarehouseManagementSystem");  
  
Container container = await database.CreateContainerIfNotExistsAsync(  
    id: "invoices",  
    partitionKeyPath: "/customerId", throughput: 400);  
  
await container.UpsertItemAsync(invoice);
```

 **Automatically serializes data to json**



There's a lot more to explore!



Redis

In-memory key value store

Distributed cache



Using Redis

```
static ConnectionMultiplexer connection =  
    ConnectionMultiplexer.Connect("CONNECTION STRING");  
  
IDatabase database = connection.GetDatabase();  
  
var cachedItem = database.StringGet(itemId);  
  
if(cachedItem.HasValue)  
{  
    return cachedItem;  
}  
  
database.StringSet(itemId, "the data");
```



Using Redis

```
static ConnectionMultiplexer connection =  
    ConnectionMultiplexer.Connect("CONNECTION STRING");
```

```
IDatabase database = connection.GetDatabase();
```

```
var cachedItem = database.StringGet(itemId);
```

```
if(cachedItem.HasValue)  
{  
    return cachedItem;  
}
```

```
database.StringSet(itemId, "the data");
```



Add caching in hot paths to
reduce load times!



Automatic Expiration of Cached Items

```
database.SetString(itemId,  
    "the data",  
    expiry: TimeSpan.FromHours(1)  
);
```



You can reference other data!

It will **not** be **enforced** by the
database.



Table Store

Built for extreme amounts of data

Very fast



Using a Table Store

```
TableServiceClient tableServiceClient = new("CONNECTION STRING");  
  
TableClient tableClient =  
    tableServiceClient.GetTableClient("SupportRequests");
```



Using a Table Store

```
TableServiceClient tableServiceClient = new("CONNECTION STRING");  
  
TableClient tableClient =  
    tableServiceClient.GetTableClient("SupportRequests");  
  
tableClient.CreateIfNotExists();
```



Using a Table Store

```
TableServiceClient tableServiceClient = new("CONNECTION STRING");

TableClient tableClient =
    tableServiceClient.GetTableClient("SupportRequests");

tableClient.CreateIfNotExists();

tableClient.AddEntity(...);

var requests = tableClient.Query<SupportRequest>(
    r => r.PartitionKey == customerId.ToString()
);
```



You now have the tools to
work with data access in your
applications!

